

# Simple TCP

EE323 Spring 2021

[ee323@nmsl.kaist.ac.kr](mailto:ee323@nmsl.kaist.ac.kr)

# Logistics

In lab sessions, we will give a brief introduction of upcoming projects.

- (4.5%) Lab session #1: Socket Programming
  - Open at 3/10, Due 3/18 (11:55 pm) - 1 week
- (4.5%) Lab session #2: HTTP Proxy
  - Open at 3/19, Due 3/30 (11:55 pm) - 1.5 week
- (6+6%) Lab session #3: Simple TCP in Reliable Transport Layer
  - Open at 3/31, Due 4/13, 5/11 (11:55 pm) - 2/4 weeks for each (including Midterm period)
- (9%) Lab session #4: Simple Router
  - Open at 5/12, Due 6/1 (11:55 pm) - 3 weeks

# The Ultimate Guide

## Primary Project Document

<http://bit.ly/ee323-proj3-1-2021>

<http://bit.ly/ee323-proj3-2-2021>

This slide is based on the document above.

Please refer to this document first if there is any question.

Still ongoing project - we need you help and participation!

You can view and comment on the document directly, so please participate.

(hey, it's a rich source of participation points!)

# STCP: a Simple Reliable Transport Layer

- Design and implement your own socket layer, MYSOCK, which supports reliable transport layer
  - Socket is a set of layers
  - You should implement only Transport layer, others are given
- STCP (Simple TCP)
  - Reliable, connection-oriented, in-order, full duplex end-to-end delivery mechanism
  - Compatible with TCP (but, it is NOT TCP)
    - No flow control, no retransmission
  - Please refer to the provided specification when in doubt

# Milestones

- Make the client/server be able to communicate on the reliable network
- Reliable mode:  
NO packet drop or out-of-order delivery in the network
- Should meet all remaining functionalities to transmit packets correctly

# Getting Started

- Read the ultimate document and RFC 793 carefully
- Download the STCP tarball from the document and extract it on one of the lab machines
  - `$ tar xzvf project3.tar.gz`
- Check any compile errors with current Makefile
  - It should compile and run on any lab machines
  - The server and client will compile; use them for testing

# Code Structure

## Application Layer

- Simple, dummy client / server
  - a. Client sends a request for a file
  - b. Server transmits a file to client
  - c. Client saves the transmitted file locally with filename “recv”
- Help with debugging of your transport layer

mysock.h

## Transport Layer: your task!

- Currently, it is just a bogus minimal transport layer
- **You should implement your own transport layer in “transport.c”**

stcp\_api.h

## Network Layer

- Emulates datagram communication service
- **Interfaces for transport layer is defined in stcp\_api.h**

# Your working playground is ...

## **transport.c**

- You will implement your transport layer in transport.c
- You are NOT allowed to modify any other .c or .h files
- You are NOT allowed to modify Makefile
- Read comments in the file carefully to understand what to do
- Consider error or corner cases and make sure to clean up dynamically-allocated memory

One thread manages only one connection.

Mysock.c calls transport\_init() to make thread.



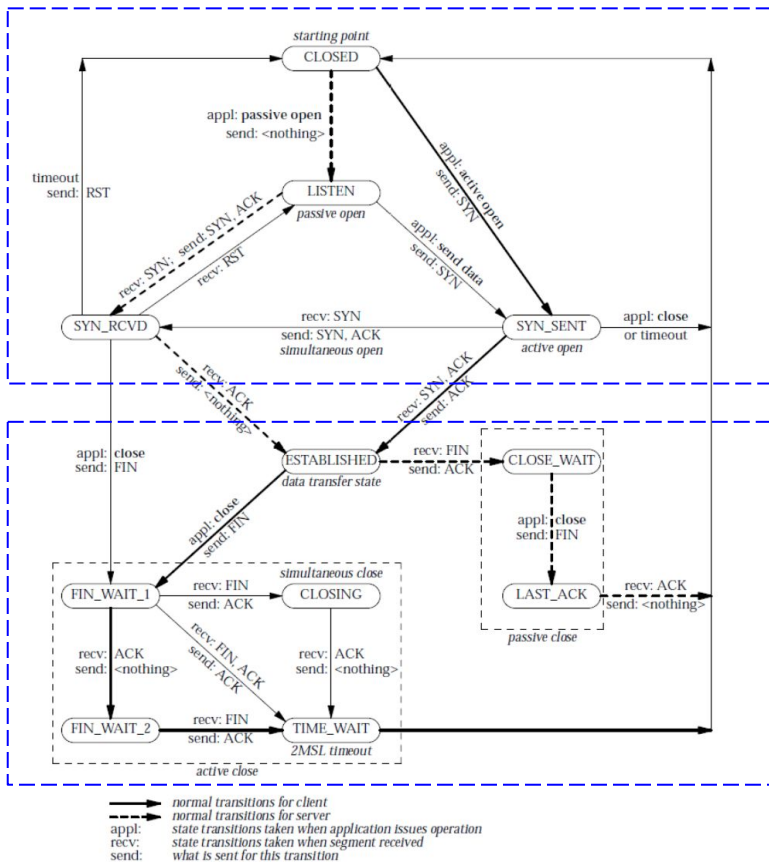
# Objectives

- In part 3-1 you will implement,
  - 3-way handshaking for connection establishment
  - 4-way handshaking for connection teardown
  - Sequence/ack number semantics
- In part 3-2 you will implement,
  - Slow start and congestion avoidance
  - Sender window management
  - Sequence/ack number semantics



## Part 3-1: TCP Handshaking

# Understand TCP Finite State Machine



Connection setup  
(3-way handshaking)

Connection teardown  
(4-way handshaking)

# Connection setup (3-way handshake)

## Active open

*client state*

LISTEN

SYNSENT

ESTAB

choose init seq num,  $x$   
send TCP SYN msg

SYNbit=1, Seq= $x$

SYNbit=1, Seq= $y$   
ACKbit=1; ACKnum= $x+1$

received SYNACK( $x$ )  
indicates server is live;  
send ACK for SYNACK;  
this segment may contain  
client-to-server data

ACKbit=1, ACKnum= $y+1$

## Passive open

*server state*

LISTEN

SYN RCVD

ESTAB

choose init seq num,  $y$   
send TCP SYNACK  
msg, acking SYN

received ACK( $y$ )  
indicates client is live

\*Borrowed from J. Kurose's slide

# Connection setup (3-way handshaking)

**`transport_init(mysocket_t sd, bool_t is_active);`**

- A connection is initialized by calling `transport_init()`
- Application function calls are blocked till the connection is established
- Make a **TCP context instance** and fill the **initial values**
- If **`is_active==TRUE`**, then your application wants to initiate a connection (e.g., called `myconnect()` at client)
  - Create and send a **SYN segment** and mark your state to **SYN\_SENT**
- If **`is_active==FALSE`**, your application is listening on a port
  - Wait for incoming data, verify a **SYN segment**, send **SYNACK**, and mark your state to **SYN\_RCVD**
- Call `control_loop()` for the main process

# Connection setup (3-way handshaking)

How do I use network layer interface for transport layer?

- Answers are in `stcp_api.c` & `stcp_api.h`
- For sending segments to peer, **`stcp_network_send()`**
- For receiving segments from peer, **`stcp_network_recv()`**
- For packing segments to be sent, refer to **`STCPHeader`** in `transport.h`
- If you want to return, and unblock the application,  
**`stcp_unblock_application()`**

Please read the code of `stcp_api.c` & `stcp_api.h` carefully!

# Connection teardown (4-way handshaking)

## Active close

*client state*

ESTAB

`clientSocket.close()`

FIN\_WAIT\_1

can no longer  
send but can  
receive data

FIN\_WAIT\_2

wait for server  
close

TIMED\_WAIT

timed wait  
for  $2 \times \text{max}$   
segment lifetime

CLOSED



FINbit=1, seq=x

ACKbit=1; ACKnum=x+1

FINbit=1, seq=y

ACKbit=1; ACKnum=y+1

can still  
send data

can no longer  
send data

## Passive close

*server state*

ESTAB

CLOSE\_WAIT

LAST\_ACK

CLOSED

\*Borrowed from J. Kurose's slide

\* For **simultaneous close**, please refer to <https://tools.ietf.org/html/rfc793> page 38-39.

# Connection teardown (4-way handshaking)

**control\_loop(mysocket\_t sd, context\_t \*ctx);**

- **Main processes** are described in this function (data transfer ~ teardown)
- Get an **event** using the return value of **stcp\_wait\_for\_event()**
  - Incoming segment from the peer (if event==NETWORK\_DATA)
  - ~~Net data from the application via mywrite() (if event==APP\_DATA) ⇒ Project 3-2~~
  - The socket to be closed via myclose() (if event==APP\_CLOSE\_REQUESTED)
- Do appropriate jobs considering the event
  - ~~Check the state, change the state, and send a packet, etc. ⇒ Project 3-2~~
  - Close the connection
- Use **TCP context** to store the state and other necessary information
- Exit the control-loop when the connection is finished



# Connection teardown (4-way handshaking)

\* In RFC793, **FINACK** is used instead of FIN. Please follow it.

How do I finish the connection in STCP api?

- API finishes the connection after calling `close()` & receiving FINACK
  - Using **`stcp_fin_received()`**, you have to notify FINACK arrival
- After the teardown, make `ctx->done` TRUE, then escape the loop
- Don't think about **TIME\_WAIT** state. Treat it as CLOSED

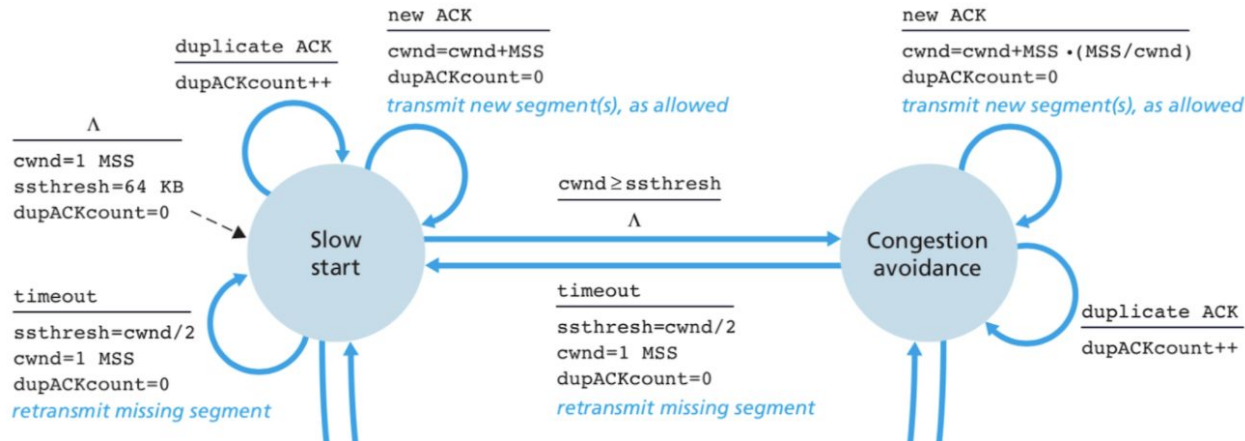
Please read the code of `stcp_api.c` & `stcp_api.h` carefully!



## Part 3-2: Reliable Data Transfer

# Reliable data transfer

- No packet loss, reordering, retransmission and timeout
- Control sender window:  $\min(\text{receiver window}, \text{congestion window})$ 
  - Receiver window size is fixed to 3072
- Implement slow start & congestion avoidance



# Reliable data transfer

**control\_loop(mysocket\_t sd, context\_t \*ctx);**

- **Main processes** are described in this function (data transfer ~ ~~teardown~~)
- Get an **event** using the return value of **stcp\_wait\_for\_event()**
  - Incoming segment from the peer (if event==NETWORK\_DATA)
  - Net data from the application via mywrite() (if event==APP\_DATA)
  - ~~— The socket to be closed via myclose() (if event==APP\_CLOSE\_REQUESTED) ⇒ Project 3-1~~
- Do appropriate jobs considering the event
  - Check the state, change the state, and send a packet, etc.
  - ~~— Close the connection ⇒ Project 3-1~~
- Use **TCP context** to store the state and other necessary information
- Exit the control-loop when the connection is finished

# Reliable data transfer

- Use **stcp\_app\_recv()** to get data from the `mywrite()` call
- Use **stcp\_app\_send()** to get data from the `myread()` call
- Please use correct *seq&ack* numbers in the transmitted segments
  - Seq number is increased **after receiving ack** for the previous segment
- Implement sender window management
  - STCP\_MSS is defined as **536** in `transport.h`
  - Start with **1 STCP\_MSS**, and the *ssthreshold* is **4 STC\_MSS**
  - In this project, assume receiver window as a fixed size of **3072**

# Testing

`client.c` and `server.c` are the codes for application calls

- Run given **client** and **server** on different shells
- Give **the name of the file** to client
  - You can use `-f [filename]` option to give the name of the file to be transferred
  - You can give the name of the file at run-time
- Client may generate a request to server and server will transmit the file as a response
- The received file at client will be saved as “**rcvd**”

Execution example)

```
$ make all
```

```
$ ./server.c
```

```
> Server's address at [myserver]:[myport]
```

```
$ ./client.c [myserver]:[myport] -f [filename]
```

# Testing

You can log the **sent/received segment information** captured from

`stcp_network_send()` and `stcp_network_recv()` function

- In **stcp\_api.c**, change the **18th line**
  - `#define LOG_PACKET FALSE` → `#define LOG_PACKET TRUE`
- Then whenever send and recv are called, logs will be written named **“pcap\_from\_[my port number]\_to\_[connected port number]”**
- Flag/sequence number/ack number/length/window size are logged
- \* It doesn't overwrite, but append the log if there is existing file (`make clean` clears it)

# Testing

- In part 3-1, we will not give file input option for the client and server
  - We will use modified client/server which calls listen/accept/connect/close in a different manner
  - You can test it by modifying client.c/server.c/transport.c
- In part 3-2, we will give file input option with -f
  - Please compare the original & transferred file with `diff`
  - We will check if there is window size increase with the log file



# Caution

- You are NOT allowed to modify or submit any other .c , .h, or Makefile in stub codes rather than 'transport.c'.
  - You will submit only 'transport.c', but not other files.
  - You can modify code for your debugging, but remember that you code should work with original Makefile and supporting code.
- Your code will be graded on haedong lounge server.
  - Make sure that your code compiles and runs properly on the machines.
- We will test correct endianness.
  - Don't forgot to include your ntohs(), htons().

# Submission

**Due date:**

**3-1 April 13th, Thursday, 23:55 PM**

**3-2 May 11th, Tuesday, 23:55 PM**

For both projects, you need to submit

- transport.c
- report.pdf

Compress above items into one zip file and rename to:  
{studentID}\_{name(in English)}\_project3\_{part#}.zip

(ex. 20219876\_JohnDoe\_project3\_2.zip)

# Grading Criteria

This project is worth 6% / 6% of your total grade.

## - Project 3-1

- (25%) Code: Active open
- (25%) Code: Passive open
- (20%) Code: Active close
- (20%) Code: Passive close
- (10%) Code: Simultaneous close
- (0%) Report

## - Project 3-2

- (40%) Code: Sender side implementation
- (40%) Code: Receiver side implementation
- (20%) Code: Congestion control implementation
- (0%) Report

- For all test cases, the order of SYN/SYNACK/ACK/FINACK flags, the corresponding sequence/ack numbers, and length should be correct. If any of them is wrong, we will score the case as 0.
- Any violations (wrong Makefile script, wrong file name, report is not PDF, etc.) will result in a penalty.

# Tips

- Print every information that you want to check the correctness
  - Most straightforward and powerful way
- Do NOT try to implement everything at once
  - Top-down implementation is important
  - Implement big branches first
  - Just describe what should be done at each block briefly
- Use dummy function that will be implemented later
  - Implement details step-by-step (test before implementing next block)
- Do NOT use global variable – Use context instance for each connection

# Others

- Do not copy and paste other's code including publicly available source code
- Start assignment **as quickly as possible**
- **Design first, before you start it**
- This assignment is newly designed, so there might be some confusing points. If you have questions, feel free to **ask question** in Campuswire